SS2P2: A Stable State-Space Point Process for Limit-Order-Book Simulation

Anonymous submission

Abstract

Simulating limit-order-book (LOB) order flow is a fundamental generative task for market-making world models, agent training, and counterfactual market analysis. Neural marked temporal point processes (MTPPs) predict the next LOB event accurately, but their simulation behaviour is largely unexamined: closed-loop, the learned self-excitation drifts or explodes, and we show that the standard remedy — rescaling the intensity to match the real event rate — often has *no solution* for unbounded models. To close this gap we propose SS2P2 (Stable State-Space Point Process), which keeps the expressive state-space backbone of S2P2 and replaces only the output heads: a *softmin-bounded rate head* whose intensity carries a hard closed-form ceiling (an exact dominating rate for Ogata thinning) with a floor of exactly zero, and a *rate-neutral* soft-max mark head that can be arbitrarily deep without perturbing the total rate. Extensive experiments on three Coinbase assets, against five baselines with three training seeds each, lossless streaming evaluation, and *verified* rate calibration for every model, show that SS2P2 is the strongest predictor on two of three assets while remaining verifiably calibratable closed-loop with competitive simulated structure. Our sharpest finding is that whether a model *can* be calibrated is itself a measurable property. The unbounded parent model fails verified calibration on 7 of 9 checkpoints — on one asset, no checkpoint can be calibrated at all — because its simulated rate responds almost discontinuously to the rescaling constant; the bounded head passes on all nine. A per-gap decomposition shows that the remaining prediction gap comes almost entirely from quiet periods, which is the expected and accepted cost of bounding the intensity. SS2P2 trains with an exact parallel scan, simulates with $O(1)$ cost per event using Ogata thinning, and matches real event rates after verification.

Introduction

The limit order book (LOB) is the core mechanism of electronic markets: a continuous-time stream of order placements, cancellations, and trades whose fine structure drives price discovery and liquidity. Two tasks are asked of a generative model of this stream. The first is *prediction*: given the event history, score the time and type of the next event. The second is *simulation*: roll the model forward freely, feeding its own samples back as history, and reproduce the stylized facts of real order flow — calibrated event rates, burst–lull clustering, and heavy-tailed returns (Cont 2001). Simulation is the harder and, for downstream use (market-making world models, agent training, counterfactual analysis), the more valuable of the two (Li et al. 2025; Nagy et al. 2025).

Neural MTPPs (Du et al. 2016; Mei and Eisner 2017; Shchur et al. 2021) excel at prediction, and continuous-time state-space variants (Chang et al. 2025; Shi and Cartlidge 2022) push held-out likelihood to the state of the art. But that literature evaluates almost exclusively teacher-forced — the S2P2 paper itself (Chang et al. 2025) reports state-of-the-art held-out likelihood yet never rolls the model out, leaving closed-loop simulation an open gap. Two obstacles stand between a trained MTPP and a usable simulator. *First, instability*: the very self-excitation that wins the one-step likelihood becomes a liability closed-loop — each sampled event raises the intensity, which produces more events, and with no mechanism bounding the rate the roll-out drifts hot or explodes. *Second, uncalibratability*: the standard patch — rescale the intensity until the simulated rate matches reality — turns out to be unavailable exactly where it is needed most. We find trained unbounded state-space checkpoints whose closed-loop rate is *near-discontinuous* in the rescaling constant (a 10^{-3} change flips the rate across the entire tolerance band) so that no constant rescaling exists — on one of our three assets, this leaves the unbounded model with no usable simulator at all. The field thus faces a trade-off between expressivity and stability: expressive models predict well but cannot be trusted (or even calibrated) in closed loop, while simple stable models simulate safely but predict poorly.

SS2P2 changes only the output heads of S2P2; the backbone is untouched. We keep the stacked latent-linear-Hawkes state-space backbone verbatim — its diagonal zero-order-hold dynamics, parallel-scan computability, and LayerNorm'd stack output h_t — and replace only what comes after h_t: a **softmin-bounded rate head**, whose intensity has a hard closed-form ceiling (an exact dominating rate for Ogata thinning) and an exactly-zero floor with unit log-intensity gradient in quiet periods; and a **rate-neutral mark head**, a soft-max over event types that cannot change the total rate no matter how deep its network is.

Because the backbone is identical, differences against the unbounded parent (S2P2) are attributable to the heads. Our contributions:

1. **The SS2P2 head pair** (softmin-bounded rate $\times$ rate-

neutral marks): it turns S2P2 — state-of-the-art teacher-forced, but unstable and often uncalibratable in free-running simulation — into a *provably stable simulator*: a bounded, thinning-exact intensity on the unmodified state-space backbone, trained with stateful (TBPTT) likelihood and simulated with $O(1)$-per-event carried state.
2. **Calibratability as a measurable model property.** Every model is rate-calibrated by a mark-preserving bisection with *full-scale verification*: the unbounded head fails on 7/9 checkpoints with a critical-point signature — including all three seeds on SOL, leaving no usable simulator — while the bounded head verifies on all nine. To our knowledge this is the first cross-asset quantification of closed-loop calibratability for neural TPPs.
3. **Exact post-hoc rate calibration** for factorized heads: the rate is linear in a learnable scale, the ceiling scales with it, and the mark distribution is untouched — a one-dimensional root-find that keeps the guarantee intact and *improves* every structure metric by correcting the simulated clock rate.

Background and Related Work

Marked temporal point processes. An MTPP models a sequence $\mathcal{H}_T = \{(t_i, e_i)\}_{i=1}^{N_T}$ of events with times t_i and types $e_i \in \mathcal{E}$, a finite set of atomic event types, through per-type conditional intensities $\lambda_e(t \mid \mathcal{H}_t)$, conditioned on the history $\mathcal{H}_t$ of all events up to time t (model parameters θ suppressed; we abbreviate $\lambda_e(t)$ when the conditioning is clear). Writing $\lambda(t) = \sum_{e\in\mathcal{E}} \lambda_e(t)$ for the total (ground) intensity and $p(e \mid t) = \lambda_e(t)/\lambda(t)$ for the mark distribution, the log-likelihood of a sequence is

$$\log \mathcal{L} = \sum_i \Big[\log \lambda(t_i) + \log p(e_i \mid t_i)\Big] - \int_0^T \lambda(u \mid \mathcal{H}_u)\, du,$$

the integral being the *compensator* $\Lambda(t) = \int_0^t \lambda(u \mid \mathcal{H}_u)\, du$. By the random time change theorem, the rescaled masses $u_i = \Lambda(t_i)$, $\Delta u_i = u_i - u_{i-1}$ are i.i.d. $\mathrm{Exp}(1)$ iff the intensity is correctly specified — we report their mean (target 1) and KS distance as timing-calibration diagnostics. Sampling uses Ogata thinning, which requires a dominating rate $\bar{\lambda} \geq \lambda(t)$; a model whose intensity has no computable upper bound must resort to heuristic (and bias-prone) dominating-rate estimates.

The S2P2 backbone. S2P2 (Chang et al. 2025; Shi and Cartlidge 2022) is a state-space point process: L stacked latent-linear-Hawkes/SSM layers with diagonal state, ZOH-discretized between events and computable by parallel scan,

$$x_{t'}^{(\ell)} = \bar{A}^{(\ell)} x_t^{(\ell)} + (\bar{A}^{(\ell)} – I)\, \tilde{B}^{(\ell)} h^{(\ell-1)}\ \big[+ \tilde{E}^{(\ell)} \alpha \text{ at an event}\big],$$

with a LayerNorm'd residual readout $h^{(\ell)} = \mathrm{LayerNorm}\big(\mathrm{GELU}(y^{(\ell)}) + h^{(\ell-1)}\big)$ producing the continuous-time hidden representation $h_t := h^{(L)}(t) \in \mathbb{R}^H$. The published model reads the rate from h_t through a projection-and-softplus head, $\lambda = \mathrm{softplus}(w^\top h_t + b)$, which is unbounded above and has no learnable scale; marks are coupled to the same rate field (our benchmark implementation of this baseline splits h_t into disjoint halves for the rate and mark heads). SS2P2 keeps everything up to and including h_t and changes only the two heads.

Evaluation axes. We score *prediction* per genuine test event (NLL splits, type accuracy, expected-time MAE, time-rescaling diagnostics) and *simulation* by calibrated closed-loop roll-outs scored on stylized facts of real order flow (Cont 2001; Vyetrenko et al. 2020); the protocol is detailed in the Experiments section.

Neural TPPs for event streams. Neural MTPPs learn history-dependent intensities with recurrent (Du et al. 2016; Mei and Eisner 2017), attention-based (Zhang et al. 2020; Zuo et al. 2020), jump-diffusion (Zhang et al. 2024), and continuous-time state-space (Chang et al. 2025) encoders (survey: Shchur et al. 2021). Shi and Cartlidge (2022) specialize neural Hawkes models to LOB event streams and evaluate both prediction and simulation. Notably, the state-space line is evaluated almost exclusively teacher-forced: Chang et al. (2025) report state-of-the-art held-out likelihood but never roll the model out (they avoid thinning altogether) — closed-loop simulation is an unaddressed gap of the S2P2 paper, and our experiments show it is not benign: the unbounded head is frequently *uncalibratable* in free roll-out (Table 3). Our contribution is orthogonal: we keep the backbone verbatim and change only what is read *out* of it.

Hawkes modeling of order flow. Parametric Hawkes processes are the classical LOB event model (Hawkes 1971, 2018; Jain et al. 2024a,b), with state-dependent extensions (Morariu-Patrichi and Pakkanen 2022). Their linear structure gives exact stationarity certificates but limited expressiveness; SS2P2 takes the complementary route — full neural expressiveness with stability restored by a closed-form bound. In the neural TPP literature such bounds appear mainly as computational devices for thinning; here the bound is a modeling decision with an asymmetric requirement (hard ceiling, exactly-zero floor), validated by our calibration study.

Generative LOB simulators. Learned order-flow generators span GANs (Li et al. 2020; Coletta et al. 2021, 2022), token-level autoregressive state-space models (Nagy et al. 2023), diffusion-guided agents (Huang et al. 2026), and foundation-model-scale transformers (Li et al. 2025; Sood et al. 2026); agent-based simulators (Byrd, Hybinette, and Balch 2020; Frey et al. 2023) remain the RL workhorse. These systems are scored on distributional realism: stylized facts (Cont 2001; Vyetrenko et al. 2020) and the KS/Wasserstein suites of LOB-Bench (Nagy et al. 2025), whose authors report error accumulation over long roll-outs. SS2P2 differs on two counts: one likelihood-based model is held to *both* prediction and closed-loop realism, and the simulated rate is calibrated and *verified* — a closed-loop stability certificate no generative simulator above carries.

Methods

SS2P2 factors the per-type intensity into a bounded total rate and a rate-neutral mark distribution — the ground-intensity–

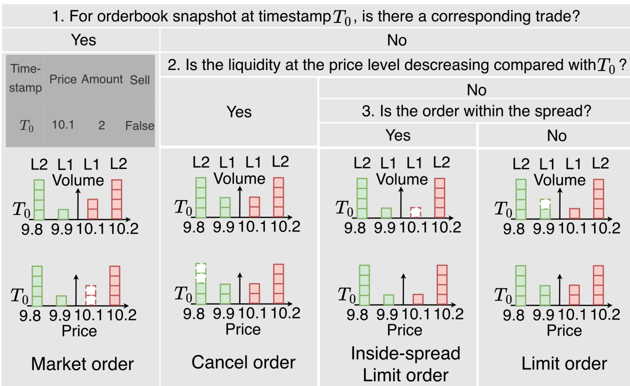

Figure 1: Event construction for the four order classes through three questions, each a decision-tree classification applied to synchronized Level-2 book snapshots and trade records at timestamps $T_0 \to T_1$: (1) is there a corresponding trade? (2) does liquidity at the price level decrease? (3) is the new order within the spread? Market order is a case of liquidity decreasing.

mark factorization of set-valued MTPPs (Chang, Boyd, and Smyth 2024) — both decoded from the *same* continuous-time hidden representation h_t:

$$\lambda_e(t) \;=\; \lambda(t) \cdot p(e \mid t), \qquad \sum\nolimits_{e \in \mathcal{E}} \lambda_e(t) = \lambda(t),$$

the second identity holding because $\sum_e p(e \mid t) = 1$. The design principle is to put nonlinearity where it cannot perturb the rate.

Event constructor

To apply point process modeling to high-frequency limit order book (LOB) data, we represent market dynamics as a sequence of **event sets** (pipeline overview in Figure 2). We consider the top $K = 10$ levels on each side (bid and ask).

Atomic Events. We define a finite set $\mathcal{E}$ of *atomic event types*, each denoted $e = (c, s, \delta)$, where:

- $c \in \{\mathrm{LO}, \mathrm{MO}, \mathrm{CO}, \mathrm{IS}\}$ is the *event class*:
  - **Limit Order (LO)**: liquidity provision at one of the top-K levels.
  - **Cancel Order (CO)**: liquidity withdrawal at one of the top-K levels.
  - **Market Order (MO)**: liquidity consumption at the best quote.
  - **Inside-Spread Order (IS)**: a new order placed inside the spread, improving the prevailing best quote.
- $s \in \{b, a\}$ is the *side* (bid or ask).
- δ is the *level index*:
  - For LO/CO, $\delta \in \{1, \ldots, K\}$ indicates distance from the best quote ($\delta = 1$ = best bid/ask, $\delta = 10$ = tenth level).
  - For MO, $\delta = 1$ always, since market orders execute against the current best quote.

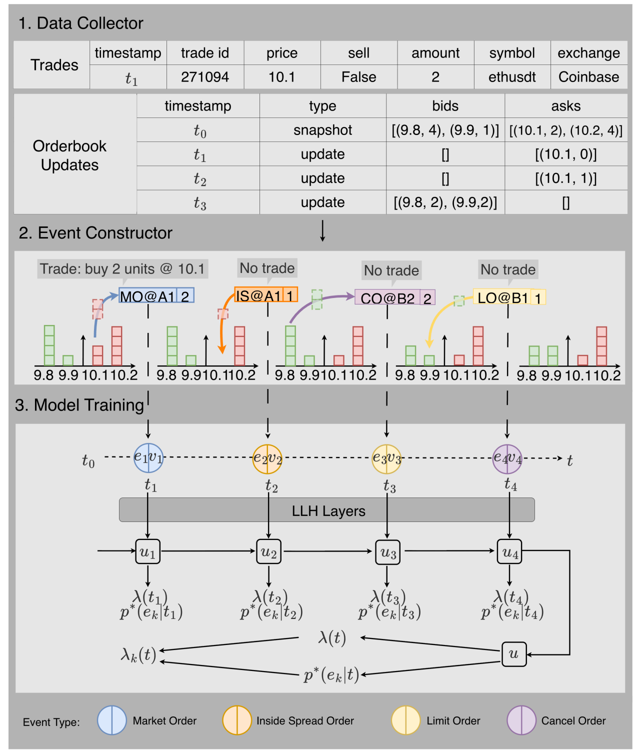

Figure 2: Pipeline. (1) *Data collector*: synchronized trade records and Level-2 order book updates. (2) *Event constructor*: consecutive book snapshots are differenced and aligned with trades to emit atomic typed events. (3) *Model training*: the event-set sequence drives the state-space backbone (LLH layers) whose continuous-time hidden representation h_t decodes the bounded total rate $\lambda(t)$ and the rate-neutral mark distribution $p(e \mid t)$, giving per-type intensities $\lambda_e(t) = \lambda(t)\, p(e \mid t)$.

  - For IS, $\delta \in \{1, \ldots, K\}$ denotes the number of ticks by which the new quote improves the spread (e.g., $\mathrm{IS}_{b,2}$ = new bid order two ticks inside the spread).

Event Construction from LOB and Trade Data. As illustrated in Figure 1, each atomic order event is inferred by differencing consecutive Level-2 order book snapshots and aligning them with trade records at identical timestamps. Because trades and book updates are perfectly synchronized in both time and volume, every observed change can be classified without ambiguity through a three-step decision tree. First, if the best-quote liquidity decreases and a trade is recorded at the same price, the change corresponds to a **Market Order**—a direct execution against standing quotes. Second, if liquidity decreases at any level without a matching trade, the reduction represents a **Cancel Order**, reflecting liquidity withdrawal. Finally, when liquidity increases, the classification depends on the price position: if the new order appears inside the spread, it forms an **Inside-Spread Order**; otherwise, it is a **Limit Order** placed within the existing book depth. This rule set ensures exhaustive mapping from raw order book and trade data into atomic order events, enabling a consistent foundation for downstream temporal modeling.

Softmin-bounded rate head

From the LayerNorm'd (hence bounded-scale) h_t, a gate produces a two-sided-bounded state g, which is read out through a smooth one-sided cap:

$$\begin{aligned} o &= \sigma(W_o h_t + b_o) \in (0,1)^H, \\ g &= o \odot \tanh(h_t) \in (-1,1)^H, \\ z_{\text{raw}} &= w^\top g + b \ \ \text{(unconstrained)}, \\ z &= z_{\max} - \operatorname{softplus}\big(z_{\max} - z_{\text{raw}}\big), \\ \lambda(t) &= \varsigma \cdot \operatorname{softplus}(z), \end{aligned} \tag{1}$$

with a learnable positive scale ς initialized so the baseline rate matches a target rate ($g \to 0 \Rightarrow \lambda \approx \varsigma \ln 2$).

Two properties follow immediately from Eq. (1):

- **Hard ceiling (simulation stability).** $z \leq z_{\max}$ always, hence $\lambda(t) \leq \varsigma \cdot \operatorname{softplus}(z_{\max})$ in closed form. This is an *exact* dominating rate for Ogata thinning — the roll-out cannot run away, whatever state the closed loop visits.
- **Open floor (prediction honesty).** As $z_{\text{raw}} \to -\infty$, $z \to z_{\text{raw}}$, so $\lambda \sim \varsigma e^{z_{\text{raw}}} \to 0$ and $\log \lambda$ keeps unit gradient in the quiet regime. The floor is exactly zero.

This asymmetry is deliberate. An earlier version of the head bounded z on *both* sides (an ℓ_1-capped readout of the gated state), giving $\lambda \in (\ell_-, \ell_+)$ with a nonzero floor ℓ_-. The ceiling ℓ_+ buys stability, but the nonzero floor hurts prediction in quiet periods: in a quiet gap of length Δt the compensator pays an irreducible log-likelihood penalty of $(\ell_- - \lambda^{\text{true}})\Delta t$ that no parameter setting can recover. The requirement is asymmetric — ceiling for simulation, zero floor for prediction — and the softmin cap builds exactly that asymmetry into the nonlinearity.

Rate-neutral mark head

Marks are decoded from the same h_t but never change the total rate:

$$p(e \mid t) = \operatorname{softmax}\big(\operatorname{MLP}(h_t)\big)_e, \qquad \lambda_e(t) = \lambda(t)\, p(e \mid t).$$

Because the soft-max lives on the simplex, the mark network can be arbitrarily deep — or later conditioned on exogenous state such as an agent's resting quotes — without touching λ. This decouples the two failure modes: rate stability is owned entirely by the bounded head, mark expressiveness entirely by the simplex.

Likelihood and sampling

Training maximizes the standard MTPP likelihood with the compensator evaluated on the training window; the mark term is a categorical cross-entropy on $p(e_i \mid t_i)$, decoupled from the rate term. Closed-loop sampling uses thinning with the exact dominating rate $\varsigma \cdot \operatorname{softplus}(z_{\max})$.

Carried-state simulation

The standard roll-out protocol re-encodes a sliding window of the last W events from a zero state at every step. This is faithful to windowed training but has two costs: $O(W)$ work per simulated event, and a hard memory ceiling — nothing older than W events can influence the future, so long-memory stylized facts (the slow decay of $|r|$ autocorrelation) are *architecturally* unreachable regardless of the model. Because the backbone is a state-space recursion, both costs vanish: we carry the packed decoder state and advance it with only the new event, which is mathematically identical to encoding the entire history from the seed, at $O(1)$ cost per event — an order of magnitude faster than window re-encoding and memory-unbounded. Attention- and window-shaped baselines have no such Markov state and retain the sliding window.

Stateful (TBPTT) training

Windowed maximum likelihood encodes every training window from a cold initial state, so the model only ever trains in a "history just began" transient and MLE explains the observed rate with an inflated baseline — the free roll-out, whose state is warm, then double-counts. We train instead with truncated backpropagation through time: batch lanes walk contiguous stretches of the stream in order, and the decoder state at the end of window t is passed — *detached* — as the initial state of window $t{+}1$ (resets only at genuine stream boundaries, to the learned initial state). Gradients remain window-local; state values flow from the start of the stream. This makes the training regime consistent with carried-state simulation, at unchanged per-step cost; it requires non-overlapping windows (stride = window length), so our controls match that setting.

Post-hoc rate calibration

The factorization buys one more property: $\lambda = \varsigma \cdot \operatorname{softplus}(z)$ is *exactly linear* in the learnable scale ς, the thinning ceiling $\varsigma \cdot \operatorname{softplus}(z_{\max})$ scales with it, and the mark head is untouched by construction. The free-running rate — which is monotone but nonlinear in ς through the closed-loop feedback — can therefore be calibrated by a one-dimensional geometric bisection on a multiplier κ: roll out briefly at $\kappa\varsigma$, compare the realized rate to the target, and iterate over a handful of probe roll-outs. The result is a certificate-preserving, sampling-exact simulator at any prescribed mean rate. For *any* decoder, a constant κ multiplying the sim-time intensity is still mark-preserving — type probabilities λ_e/λ are unchanged — so the same bisection applies to every baseline; what is lost off-family is only the exactness of the thinning bound. We make the procedure falsifiable: the target is the *validation*-split rate, the bisection must bracket and converge, and a full-scale validation roll-out must verify the target within tolerance (failures escalate to full-scale probes; unresolved failures exclude the checkpoint and are reported). Transfer to the test split is reported, never gated. We use calibration as a *post-hoc* correction for the residual rate inflation caused by the biased endpoint-rule training compensator — and, in the experiments, as a measurement instrument in its own right.

Experiments and Evaluations

Protocol

We benchmark on Coinbase BTC, ETH, and SOL event-driven order flow, seven days per asset (validation-split rates 38.3, 47.9, and 24.0 events/s), encoded by the TFOW event constructor (Anonymous 2026) of the Methods section: $|\mathcal{E}| = 62$ atomic types ($20+20+2+20$ over LO/CO/MO/IS, sides, and levels). Six models — NHP (the Mei–Eisner CT-LSTM), an LSTM decoder and a SAHP-style causal-attention decoder (Zhang et al. 2020) (both *adapted generic backbones* over our shared input embedding, not paper-faithful re-implementations), PCT-LSTM (a per-type parallel CT-LSTM), S2P2, and SS2P2 — are trained under an identical configuration per asset: sequence length 1024, non-overlapping windows, identical gradient-step counts, three training seeds each. The fixed 1024-event context spans only $\sim$20–40 s of market time at these rates, so long-memory structure leans entirely on carried state. The protocol is strict end-to-end:

- **Prediction** is scored on *every* genuine test event by a lossless streaming evaluator: raw test segments are traversed in stream order with recurrent state carried across chunks (one cold start per genuine segment boundary); the evaluator hard-fails if any event goes unscored.
- **Simulation** uses carried-state closed-loop roll-outs (32×600 s, three roll-out seeds per checkpoint) compared to equal-duration bootstrap samples of real segments. SS2P2 samples by Ogata thinning with its exact closed-form ceiling; baselines, which admit no exact bound, share a common compensator-inversion sampler.
- **Every model is rate-calibrated** before scoring: a one-dimensional bisection on a time-rescaling constant κ (mark-preserving for every decoder; see Method) targets the *validation*-split rate, then a full-scale validation roll-out must verify the target within 15%; failures escalate to full-scale probes and, if still unmet, the checkpoint is excluded and the failure reported as a result. SAHP is the exception: its closed-loop rate is unstable across warm-starts and no constant κ transfers, so it is reported uncalibrated (k=1, marked †).
- **Statistics**: roll-out seeds are averaged within each checkpoint; 95% CIs are t-based across checkpoints. Any batch failure or non-finite value aborts the run.

Prediction, and where its likelihood is lost

Table 1: SS2P2 is the strongest predictor on two of three assets — it wins *every* prediction metric on SOL by a wide margin (0.27 vs. 0.53 overall NLL; 0.192 vs. 0.184 accuracy), wins overall NLL and accuracy outright on ETH (-0.87 vs. NHP's -0.81; 0.345 vs. 0.341), and on BTC ties NHP on marks while conceding overall likelihood — all while remaining verifiably calibratable with competitive simulated structure throughout (Tables 2 and 3). The gap is in timing. A per-gap decomposition of timeNLL $= -\log\lambda(t_i) + \Delta u_i$ (event term + compensator mass) shows exactly where it comes from on BTC, the asset where NHP leads. Burst sharpness is *not* the issue: the state-space family's event terms are better than NHP's (-4.95 SS2P2 vs. -4.23), and intensity-ceiling saturation is comparable across models ($\sim$25%). The entire gap is compensator mass in *quiet gaps*: mean per-gap rescaled mass is 1.88 (SS2P2) and 2.26 (S2P2) against NHP's near-ideal 1.03 (target 1), with the excess accruing in gaps beyond 0.1 s — rare at 38 ev/s, but expensive (SS2P2 integrates 10.0 units of hazard across a 0.5–2 s gap where NHP integrates 6.0; S2P2 16.4). Figure 3 shows why: NHP's per-dimension decay toward a *learned steady state* tracks the falling empirical hazard of real gaps, while the state-space family — whose quiet-state behaviour is an emergent readout of decaying latents with no dedicated resting-rate parameter — stays hot through quiet gaps. This cost is the deliberate price of stability; what it buys is the subject of the next two subsections.

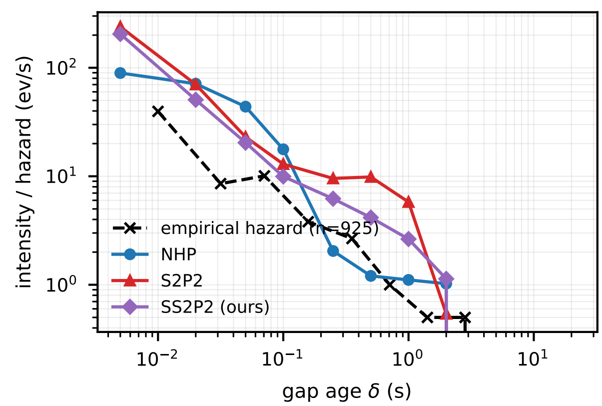

Figure 3: Mean model intensity at gap age δ, among BTC test gaps that survive past δ, vs. the empirical hazard of the real gap distribution (log-log; survivor counts shrink toward the right). NHP tracks the falling hazard; the state-space family stays hot through quiet gaps — the timeNLL gap of Table 1 accrues here.

Simulation

Table 2 and Figure 6 score calibrated closed-loop roll-outs. Every calibrated model lands its rate (rel-err 0.06–0.42); the discriminating columns are structural. PCT-LSTM leads the Fano column on all three assets; clustering splits between NHP (BTC, ETH) and SS2P2 (SOL); SS2P2 is competitive on every structural column of every asset — the only model in the table that combines that with top-tier prediction — while the LSTM sampler pins Fano rel-err at $\approx$0.99 throughout: it simulates a near-Poisson stream, the trivially-stable end of the frontier. Figure 4 shows the level-and-rise view on BTC: real Fano rises from $\sim$60 at 1 s to $\sim$1000 at 50 s; every model under-disperses at this context length, PCT-LSTM and SS2P2 least. The one surviving S2P2 arm on ETH calibrates in rate but over-disperses wildly (Fano rel-err 6.0) — rate calibration cannot repair a near-critical checkpoint's structure. NHP, the best BTC predictor, is mid-pack in simulation on every asset: prediction and simulation quality remain distinct axes, and no unbounded model wins both.

model	BTC		ETH		SOL	
	overall↓	ACC↑	overall↓	ACC↑	overall↓	ACC↑
NHP	**-1.07±0.01**	**0.449±0.003**	-0.81±0.03	0.341±0.006	0.53±0.01	0.184±0.004
LSTM	-0.44±0.40	0.411±0.015	0.15±0.04	0.288±0.003	1.34±0.25	0.141±0.005
SAHP †	-0.93±0.01	0.412±0.001	-0.60±0.01	0.297±0.004	0.67±0.02	0.155±0.003
PCT-LSTM	-0.60±0.08	0.313±0.010	-0.34±0.10	0.248±0.007	0.89±0.03	0.105±0.004
S2P2	-0.39±0.25	0.439±0.002	-0.38±0.22	0.331±0.008	0.67±0.49	0.182±0.002
SS2P2	-0.87±0.12	0.448±0.001	**-0.87±0.08**	**0.345±0.003**	**0.27±0.05**	**0.192±0.001**

Table 1: Prediction per genuine test event (mean ± 95% CI over three training seeds; streaming lossless evaluator; overall = time + mark NLL). SS2P2 wins both columns on ETH and SOL and ties NHP on BTC accuracy (0.447 vs. 0.449). Not shown: S2P2's expected-time MAE is pathological on every asset (1.5, 10.2, 46.2 s vs. 0.03–0.05 s for every bounded or recurrent baseline) — the intensity-collapse tail of the unbounded head.

model	BTC			ETH			SOL		
	rate_re	Fano_re	clus_re	rate_re	Fano_re	clus_re	rate_re	Fano_re	clus_re
NHP	0.09±0.08	0.86±0.07	**0.27±0.04**	0.37±0.20	0.87±0.11	**0.52±0.23**	**0.19±0.03**	0.75±0.17	0.72±0.11
LSTM	0.10±0.09	0.99±0.00	0.50±0.41	0.40±0.03	0.99±0.00	0.77±0.29	0.23±0.04	0.99±0.00	0.97±0.03
SAHP †	0.71±0.38	0.77±0.19	0.93±2.49	0.79±0.12	0.82±0.23	0.34±0.73	0.68±0.15	0.92±0.08	0.65±0.25
PCT-LSTM	**0.06±0.06**	**0.79±0.12**	0.36±0.30	0.42±0.26	**0.76±0.17**	0.55±0.50	0.22±0.25	**0.46±0.34**	0.72±1.35
S2P2	0.32±1.51	1.25±12.45	0.29±0.14	**0.27**	6.05	0.60	–	–	–
SS2P2	0.25±0.43	0.81±0.37	0.29±0.17	0.42±0.17	0.79±0.05	0.81±0.11	0.30±0.17	0.60±0.60	**0.68±0.62**

Table 2: Closed-loop stylized facts (rel-err vs. real; roll-out seeds averaged per checkpoint, mean ± 95% CI over checkpoints). Checkpoints that failed calibration verification are excluded (BTC: S2P2-s1/s2, SS2P2-s3; ETH: S2P2-s1/s2; SOL: all three S2P2 seeds — "–"). † = SAHP uncalibrated by protocol; its free-running rate collapses to 6–10 ev/s against 24–48 real.

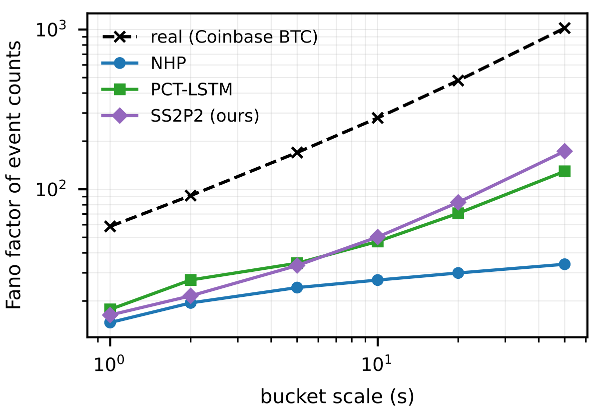

Figure 4: Fano factor of simulated event counts vs. bucket scale (Coinbase BTC; mean over calibrated checkpoints and roll-out seeds). Real order flow is super-Poissonian with Fano *rising* in scale; all models under-disperse at this context length, SS2P2 and PCT-LSTM least.

model	BTC 38 ev/s	ETH 48 ev/s	SOL 24 ev/s
NHP	3/3	3/3	3/3
LSTM	3/3	3/3	3/3
SAHP †	*model-level divergence: reported uncalibrated*		
PCT-LSTM	3/3	3/3	3/3
S2P2	**2/3**	**1/3**	**0/3**
SS2P2	3/3	3/3	3/3

Table 3: Checkpoints passing calibration + full-scale verification, out of three training seeds per asset (bold marks failures). The unbounded S2P2 head fails on 7 of 9 checkpoints — *all three* on SOL, leaving no usable simulator — while the bounded SS2P2 head fails once in nine. SAHP diverges at the model level on every asset.

Calibratability is a model property

Post-hoc rate calibration (Method) is mark-preserving for *every* decoder in the zoo — a constant κ rescales sim-time intensity without touching type probabilities — so we calibrate every model, and treat the *outcome* of calibration as an experimental result in its own right. The protocol is deliberately falsifiable: probes must bracket and converge to the validation target, and a full-scale validation roll-out must verify it; every failure is reproducible (probe seeds are deterministic, and each retried failure reproduced its bisection bracket bit-identically).

Table 3 is, to our knowledge, the first cross-asset quantification of closed-loop calibratability for neural TPPs: at production event rates, the unbounded head is uncalibratable on 7 of 9 checkpoints, and on SOL the model class is excluded outright. Figure 5 shows the failure signature: the closed-loop rate response of a failing S2P2 checkpoint is near-discontinuous in κ — a 10^{-3} change in κ flips the realized rate across the entire tolerance band (e.g. 36 → 48.5 ev/s), and seed-to-seed variance at fixed κ reaches $\sim$30% — the

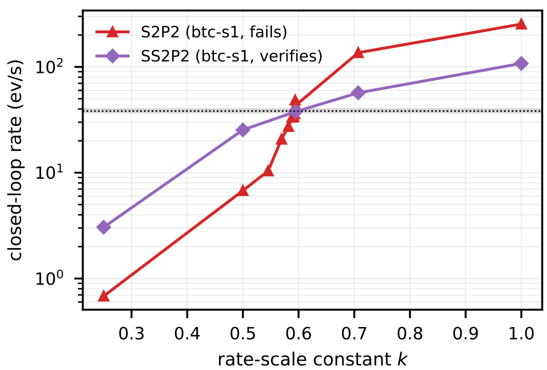

Figure 5: Calibration probe ladders on BTC (target band shaded): the failing S2P2 checkpoint's closed-loop rate is near-discontinuous in the rescaling constant κ; the verified SS2P2 checkpoint crosses the target smoothly.

behaviour of a system parked near criticality, where no constant time-rescaling exists. The bounded SS2P2 response crosses the same target smoothly and verifies on all nine checkpoints. (An earlier training run of one BTC seed produced a checkpoint whose single failure was qualitatively different — a per-lane diagnostic traced it to a *metastable* burst-locked regime at $\sim 2\times$ the target rate that individual roll-out lanes enter stochastically, bounded and benign next to S2P2's criticality; retraining escaped that basin.) This confirms the stability certificate empirically: the same unbounded softplus head whose intensity-collapse states produce S2P2's pathological expected-time tail (Table 1) also places checkpoints near criticality in closed loop, and both problems disappear under the softmin bound.

Training and simulation efficiency

The state-space recursion is not just statistically but computationally advantageous — once its per-event sequential loop is replaced by an exact associative prefix scan (Hillis–Steele over the ZOH recurrence; identical states and gradients to 2×10^{-4}). Table 4 reports wall-clock measurements on one GPU with the benchmark configuration and ETH checkpoints: one training step (fwd+bwd, batch 64×1024 events) and closed-loop simulation throughput (simulated events per wall-second, 8×60 s carried-state roll-outs at native rates; SAHP re-encodes a sliding window). The scan collapses the S2P2-family training step from 3.5 s to ~ 0.06 s — attention-class speed, $60\times$ faster than the loop and $73\times$ faster than the CT-LSTM — turning a multi-hour training run into minutes. For sampling, SS2P2's exact ceiling makes Ogata thinning available with a *constant* dominating rate (no per-step bound estimation; mean acceptance 0.54 on ETH): exactness costs $\sim 2.5\times$ throughput against the approximate quadrature-inversion sampler shared by all models, and both run far above real time (353 vs. 48 events/s on the fastest asset).

	train step (s)	vs. seq. loop	sim. ev/s
NHP	4.35	–	1253
LSTM	0.019	–	2083
SAHP	0.059	–	583
PCT-LSTM	0.79	–	1919
S2P2	0.055	$64\times$	1163
SS2P2	0.059	$60\times$	879 / 353†

Table 4: Efficiency on one GPU (ETH checkpoints, benchmark configuration). Training: one fwd+bwd on 64×1024 events (scan for the S2P2 family; speedup vs. their sequential loop). Simulation: events per wall-second, 8×60 s roll-outs. †inversion / exact-ceiling Ogata thinning.

Discussion and Conclusion

We held six neural MTPPs to both sides of the prediction–simulation tension on three Coinbase assets at production event rates (24–48 events/s), under a protocol designed so that failures are visible and reported: lossless streaming prediction, multi-seed checkpoint-level confidence intervals, and verified rate calibration whose failures count as results. The SS2P2 heads — a softmin-bounded rate with a closed-form thinning ceiling and an exactly-zero floor, times a rate-neutral simplex mark head on an unchanged state-space backbone — win every prediction metric on SOL, lead ETH, tie the strongest baseline on BTC marks, and remain verifiably calibratable with competitive simulated structure throughout, while a per-gap hazard decomposition shows that the remaining prediction gap comes from quiet periods — the expected cost of bounding the rate. The sharpest finding is that *calibratability is a model property*: the unbounded parent head fails verified rate calibration on 7 of 9 checkpoints — all three on SOL — with a critical-point signature, while the bounded head verifies on all nine. This confirms the stability certificate empirically. Stateful training and carried-state roll-out contribute temporal structure that calibration alone cannot. Because the mark head is rate-neutral, conditioning it on exogenous state — an agent's resting quotes — cannot destabilize the rate, making SS2P2 a suitable environment model for market-making world models; that action-conditioned extension is future work. The same factorized intensity also supports the importance-sampling probabilistic queries of Chang, Boyd, and Smyth (2024) — "event A before event B, given the history" — which map directly onto the quantities a market maker needs from a world model: fill probability (an aggressing order reaches our resting level before the quote moves away) and adverse-selection risk (the probability that flow continues against us conditional on a fill), estimated consistently from one likelihood-trained model rather than by separate ad-hoc regressions. All tables and figures are regenerated end-to-end by the released collection and display scripts.

References

Anonymous. 2026. TFOW: Detecting Anomalous Order Flow in Crypto Limit Order Books. Companion manuscript,

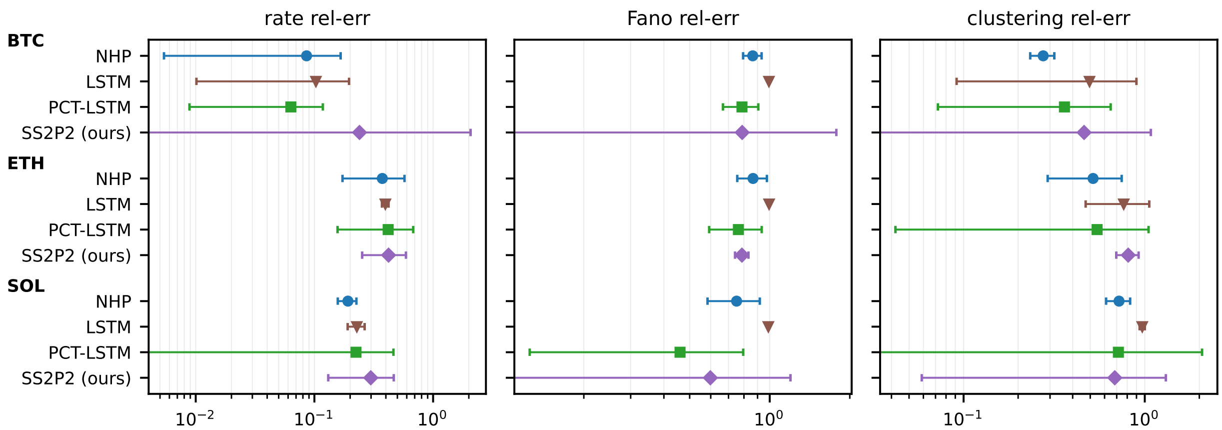

Figure 6: Simulation quality across assets (rel-err vs. real, log scale; whiskers = 95% CI over checkpoints). S2P2 and SAHP omitted where no calibrated checkpoints exist.

under review.

Byrd, D.; Hybinette, M.; and Balch, T. H. 2020. ABIDES: Towards High-Fidelity Multi-Agent Market Simulation. In *Proceedings of the 2020 ACM SIGSIM Conference on Principles of Advanced Discrete Simulation*, 11–22.

Chang, Y.; Boyd, A.; and Smyth, P. 2024. Probabilistic Modeling for Sequences of Sets in Continuous-Time. In *Proceedings of the 27th International Conference on Artificial Intelligence and Statistics (AISTATS)*, volume 238 of *PMLR*.

Chang, Y.; Boyd, A.; Xiao, C.; Kass-Hout, T.; Bhatia, P.; Smyth, P.; and Warrington, A. 2025. Deep Continuous-Time State-Space Models for Marked Event Sequences. In *Advances in Neural Information Processing Systems 38*. ArXiv:2412.19634.

Coletta, A.; Moulin, A.; Vyetrenko, S.; and Balch, T. 2022. Learning to Simulate Realistic Limit Order Book Markets from Data as a World Agent. In *Proceedings of the 3rd ACM International Conference on AI in Finance*.

Coletta, A.; Prata, M.; Conti, M.; Mercanti, E.; Bartolini, N.; Moulin, A.; Vyetrenko, S.; and Balch, T. 2021. Towards Realistic Market Simulations: A Generative Adversarial Networks Approach. In *Proceedings of the Second ACM International Conference on AI in Finance*.

Cont, R. 2001. Empirical Properties of Asset Returns: Stylized Facts and Statistical Issues. *Quantitative Finance*, 1(2): 223–236.

Du, N.; Dai, H.; Trivedi, R.; Upadhyay, U.; Gomez-Rodriguez, M.; and Song, L. 2016. Recurrent Marked Temporal Point Processes: Embedding Event History to Vector. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 1555–1564.

Frey, S.; Li, K.; Nagy, P.; Sapora, S.; Lu, C.; Zohren, S.; Foerster, J.; and Calinescu, A. 2023. JAX-LOB: A GPU-Accelerated Limit Order Book Simulator to Unlock Large-Scale Reinforcement Learning for Trading. In *Proceedings of the 4th ACM International Conference on AI in Finance*.

Hawkes, A. G. 1971. Spectra of Some Self-Exciting and Mutually Exciting Point Processes. *Biometrika*, 58(1): 83–90.

Hawkes, A. G. 2018. Hawkes processes and their applications to finance: a review. *Quantitative Finance*, 18(2): 193–198.

Huang, Y.-H.; Xu, C.; Liu, Y.; Liu, W.; Li, W.-J.; and Bian, J. 2026. Controllable Financial Market Generation with Diffusion Guided Meta Agent. In *Proceedings of the 40th AAAI Conference on Artificial Intelligence*.

Jain, K.; Firoozye, N.; Kochems, J.; and Treleaven, P. 2024a. Limit Order Book dynamics and order size modelling using Compound Hawkes Process. *Finance Research Letters*, 69: 106157.

Jain, K.; Firoozye, N.; Kochems, J.; and Treleaven, P. 2024b. Limit Order Book Simulations: A Review. *arXiv preprint arXiv:2402.17359*.

Li, J.; Liu, Y.; Liu, W.; Fang, S.; Wang, L.; Xu, C.; and Bian, J. 2025. MarS: A Financial Market Simulation Engine Powered by Generative Foundation Model. In *International Conference on Learning Representations*.

Li, J.; Wang, X.; Lin, Y.; Sinha, A.; and Wellman, M. P. 2020. Generating Realistic Stock Market Order Streams. In *Proceedings of the 34th AAAI Conference on Artificial Intelligence*, 727–734.

Mei, H.; and Eisner, J. M. 2017. The Neural Hawkes Process: A Neurally Self-Modulating Multivariate Point Process. In Guyon, I.; Luxburg, U. V.; Bengio, S.; Wallach, H.; Fergus, R.; Vishwanathan, S.; and Garnett, R., eds., *Advances in Neural Information Processing Systems*, volume 30. Curran Associates, Inc.

Morariu-Patrichi, M.; and Pakkanen, M. S. 2022. State-dependent Hawkes processes and their application to limit order book modelling. *Quantitative Finance*, 22(3): 563–583.

Nagy, P.; Frey, S.; Li, K.; Sarkar, B.; Vyetrenko, S.; Zohren, S.; Calinescu, A.; and Foerster, J. 2025. LOB-Bench:

Benchmarking Generative AI for Finance – an Application to Limit Order Book Data. In *Proceedings of the 42nd International Conference on Machine Learning*.

Nagy, P.; Frey, S.; Sapora, S.; Li, K.; Calinescu, A.; Zohren, S.; and Foerster, J. 2023. Generative AI for End-to-End Limit Order Book Modelling: A Token-Level Autoregressive Generative Model of Message Flow Using a Deep State Space Network. In *Proceedings of the 4th ACM International Conference on AI in Finance*.

Shchur, O.; Türkmen, A. C.; Januschowski, T.; and Günnemann, S. 2021. Neural Temporal Point Processes: A Review. *IJCAI International Joint Conference on Artificial Intelligence*, 4585–4593.

Shi, Z.; and Cartlidge, J. 2022. State Dependent Parallel Neural Hawkes Process for Limit Order Book Event Stream Prediction and Simulation. In *Proceedings of the 28th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, 1607–1615. New York, NY, USA: ACM. ISBN 9781450393850.

Sood, S.; Kawawa-Beaudan, J.; Borrajo, D.; and Veloso, M. 2026. TradeFM: A Generative Foundation Model for Trade-flow and Market Microstructure. *arXiv preprint arXiv:2602.23784*.

Vyetrenko, S.; Byrd, D.; Petosa, N.; Mahfouz, M.; Dervovic, D.; Veloso, M.; and Balch, T. 2020. Get Real: Realism Metrics for Robust Limit Order Book Market Simulations. In *Proceedings of the First ACM International Conference on AI in Finance*.

Zhang, Q.; Lipani, A.; Kirnap, O.; and Yilmaz, E. 2020. Self-Attentive Hawkes Process. In *International Conference on Machine Learning*, 11183–11193.

Zhang, S.; Zhou, C.; Liu, Y.; Zhang, P.; Lin, X.; and Ma, Z.-M. 2024. Neural Jump-Diffusion Temporal Point Processes. In *Proceedings of the 41st International Conference on Machine Learning*.

Zuo, S.; Jiang, H.; Li, Z.; Zhao, T.; and Zha, H. 2020. Transformer Hawkes Process. In *Proceedings of the 37th International Conference on Machine Learning*.