

TFOW: Detecting Anomalous Order Flow in Crypto Limit Order Book: [Experiments & Analysis]

ABSTRACT

The limit order book (LOB) is the core mechanism of crypto centralized exchanges by continuously recording bids and asks to facilitate price discovery and maintain liquidity, yet it is also the primary channel through which most forms of cryptocurrency market manipulation are executed. Prior studies have examined pump-and-dump schemes and wash tradings which undermine the integrity of the market, but these approaches operate at the token or exchange level, leaving open the critical question of which specific orders or order flows are manipulative. Without this granularity, regulators and exchanges lack actionable countermeasures against specific manipulators. To address this gap, we introduce the TFOW framework, which constructs atomic order events from synchronized trade and order book update data, and models their dynamics with a set-valued marked temporal point process (set-valued marked temporal point process (set-mtpp)). Leveraging the learned intensity function, we rescale inter-arrival times and apply the sum-of-squared-spacings (3S) statistic for principled anomaly detection at the sequence level. Detected anomalies are then localized through set-valued marked temporal point process (set-mtpp)-based probabilistic queries, which provide timestamp-level resolution and surface rare temporal motifs characteristic of manipulative behavior.

Our Tempo Focused Order Watcher (TFOW) [PLACEHOLDER FOR RESULTS].

Overall, Tempo Focused Order Watcher (TFOW) delivers interpretable order-level anomaly detection and equips regulators and exchanges with actionable tools for market integrity, offering a generalizable framework for identifying anomalous event flows in continuous time.

ACM Reference Format:

. 2026. TFOW: Detecting Anomalous Order Flow in Crypto Limit Order Book: [Experiments & Analysis]. In . ACM, New York, NY, USA, 9 pages. <https://doi.org/10.1145/nnnnnnnn.nnnnnnnn>

1 INTRODUCTION

The majority of crypto trading occurs on centralized exchanges (CEXs), with decentralized exchanges (DEXs) accounting for only about 10–25% of spot trading volumes, making CEX the primary venues where market manipulations such as pump-and-dump and wash trading take place.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.
Conference'17, July 2017, Washington, DC, USA

© 2026 Copyright held by the owner/author(s). Publication rights licensed to ACM.
ACM ISBN 978-x-xxxx-xxxx-x/YY/MM
<https://doi.org/10.1145/nnnnnnnn.nnnnnnnn>

In crypto CEXs, transactions are facilitated through a limit order book, which continuously maintains two sorted queues of quoted prices and quantities—bids for purchase orders and asks for sell orders [5]. To influence market dynamics, *manipulators* strategically inject orders into the limit order book, distorting the natural order flows and jeopardizing the integrity of the crypto market. As a result of these malicious orders, the order flow becomes distorted, leading to inefficient price discovery or artificially inflated trading volumes.

Previous studies have made efforts to detect such manipulations from various macro perspectives. For instance, Xu [14] investigates **pump-and-dump** schemes and demonstrate the feasibility of predicting coordinated price surges in advance. However, their framework remains token-level and does not identify which specific orders or trading actions drive these movements. Likewise, Cong [3] analyze **wash trading** using aggregate statistical techniques, providing insights at the exchange or token level. Yet, this coarse granularity overlooks the microstructural dynamics within the limit order book, where manipulative behaviors are actually executed. Understanding manipulation at this granular level—where order submissions, cancellations, and executions interact—is crucial for uncovering how artificial flows distort market.

This study presents Tempo Focused Order Watcher (TFOW), a system for unveiling manipulative behaviors within CEX limit order books. TFOW jointly detects anomalous order flows, order motifs, and individual orders under a single framework. We collect trade and Level-2 order book data, which capture aggregated market depth without individual order IDs, and reconstruct atomic order events. Building on this representation, TFOW employs sum-of-squared-spacings statistics and set-valued marked temporal point process (set-mtpp)-based anomaly queries to identify deviations from normal market dynamics.

To evaluate TFOW,

Our work makes the following contributions to the modeling and detection of anomalies in high-frequency limit order book (LOB) data:

i) **Volume-Set MTPP for LOB Dynamics.** We introduce the first volume-set marked temporal point process (volume-set MTPP) for modeling order book event dynamics. Unlike prior approaches that either discard simultaneous events or collapse them into a single representative, our formulation preserves the full event set with associated volumes, achieving a state-of-the-art generative model of microstructural dynamics.

ii) **Anomaly Detection via 3S and Query Localization.** We are the first to apply sum-of-squared-spacings (3S) statistics to rescaled sequences to detect anomalous subsequences. To further localize anomalies, we design volume-set MTPP-based probabilistic queries that surfaces rare temporal motifs underlying manipulative behaviors.

iii) **Order-Level Granularity for Regulatory Auditing.** !!!!!!!our findings Our detector operates directly at the order and trade level.

By indexing every atomic event, we can identify which specific orders deviate from the model distribution. This provides regulators and centralized exchanges with a practical screening tool for auditing of orderbooks, enabling detection of manipulative activities such as spoofing and crowd-pumps with actionable granularity.

iv) **Generalizable Anomaly Detection Framework.** Beyond order book anomaly detection, our framework extends naturally to detecting anomalies in general continuous-time event data, where events can be indexed, modeled using the set-mtpp, and evaluated to reveal statistically irregular or unexpected temporal patterns.

2 BACKGROUND

2.1 Temporal Point Processes and Order Book Modeling

Temporal point process (TPP) model event arrivals in continuous time via a conditional intensity $\lambda^*(t)$ [11]. Marked TPPs (MTPPs) [13] enrich this framework by attaching categorical marks to each event, capturing both timing and type. In financial markets, MTPPs are widely applied to model limit order book dynamics, where marks naturally encode order types (e.g., limit, cancel, market). A prominent example is the Hawkes process, a self-exciting MTPP extensively used for order book modeling due to its ability to capture clustering in order arrivals [6, 7]. The neural Hawkes process [9] extends this idea with neural architectures that learn non-linear, history-dependent intensities, outperforming classical parametric forms. However, these models assume a single event per timestamp, an unrealistic constraint in limit order books where multiple actions often co-occur. To overcome this, Chang et al. [1] introduce the set-valued MTPP (Set-MTPP), which generalizes marks from singletons to sets, enabling principled modeling of simultaneous market events. This foundation is critical for capturing realistic order flow and underpins our anomaly detection framework.

2.2 Market Participants and Manipulators

In centralized exchanges and under normal trading conditions, order flows predominantly arise from two types of participants: investors, who submit directional orders to express market views, and market makers, who continuously quote on both sides of the book to provide liquidity while managing inventory risk [4]. Under this natural market composition, the aggregate order flow exhibits a relatively stable probabilistic structure. However, when **manipulators** intervene—through practices such as pump-and-dump schemes or wash trading—they deliberately inject strategic orders into this equilibrium, distorting the underlying dynamics generated by investors and market makers and undermining the integrity of the market.

2.3 Anomaly Detection

1. deviation from the model. First order moment, second order moment 2. Querying

3 EMPIRICAL - INTO SYSTEM DESIGN

3.1 Data Selection

1. heatmap dynamic pump vs non pump

We process the full orderbook data events but keep only the top 10 for orderbook modeling, which is the practice used in the industry [12].

Memory issue when processing the orderbook construction

Since USDT pairs are not available across all exchanges, when there is not USDT as the base commodity we choose the USD as the base, since the USDT is pegged to USD and the exchange rate stay 1 to 1 ratio normally.

The order book data from Kaiko are records of updates or snapshots per timestamps, but there are rare errors where duplicate timestamps exist—less than 0.1% of the overall orderbook snapshots. In this case, we save the snapshots of both duplicate timestamps but skip the event generations

According to [3] Coinbase, Gemini for regulated, normal Binance, Huobi for unregulated, normal Bitfinex, Kucoin for unregulated, abnormal

BTC highest rejection XRP second highest rejection ETH lowest reject

Most of the exchanges have multiple update per exchange, although Gemini has only one update per timestamp.

The Kaiko steam has "u" and "s" in the order book data steam, which stand for updates and snapshots. Most of the exchanges has dominate updates and a few snapshots. However, for Kucoin, all of its steams are snapshots.

- TOP MARKET CAP (6): BTC, ETH, SOL, XRP, TON, AVAX
- MAJOR L1/L2 BLOCKCHAINS (5): DOT, POL, INJ, XTZ, FET
- DEFI PROTOCOLS (8): AAVE, UNI, COMP, CRV, LDO, SNX, SUSHI, YFI
- MEME COINS (6): DOGE, SHIB, PEPE, FLOKI, BONK, WIF
- GAMING/METaverse (3): SAND, AXS, APE
- ORACLE/DATA INFRASTRUCTURE (4): LINK, GRT, FIL, AR
- PRIVACY/PAYMENTS (2): ZEC, LTC
- DEX/TRADING INFRASTRUCTURE (3): LRC, ZRX, SKY
- OTHER/UTILITY (2): AMP, CHZ

1. how the orderbook looks like in statistics and whats the abnormal

2. pump-and-dump period order frequency
3. wash trading exchange level order amount
4. spread big or small

4 SYSTEM DESIGN

!!The state transition matrix $\phi[K, W, W]$ is trained offline via simple MLE (counting transitions), completely separate from the neural network.

We illustrate the pipeline of TFOV in Figure 1, which is composed of three key modules: (i) data collector; (ii) event constructor; (iii) anomaly detector. The data collector gathers trades and limit order book updates, which are then transformed by the event constructor into a sequence of order events. These sequences serve as inputs to the anomaly detector, which first identifies abnormal sequences using the sum-of-squared-spacings statistic, and subsequently localizes the anomalies via hitting-time and A-before-B query queries. Note, the anomaly detector starts by raising deviation alerts at the sequence level, and then pinpoints anomalies at timestamps and uncovers the underlying temporal motifs.




Figure 1: TFOW consists of three modules: (i) data collector; (ii) Event constructor; and (iii) anomaly detector. Note, (a, b) here is not an event but the actual price and depth.

4.1 Data Collector

We obtain both trade and Level-2 limit order book data from Kaiko [8], a leading institutional provider of digital asset market data. In partnership with Cloudburst Technologies [2], a leading firm specializing in detecting and mitigating market manipulation, we establish the data collector. Level-2 limit order book updates are continuously streamed to a Google Cloud Storage bucket, providing full-depth market coverage. Each update record contains precise timestamps, update types (snapshots or liquidity updates), and the corresponding price-volume information for bid and ask levels. Complementary trade data are retrieved via Kaiko’s API and temporally aligned with the order book stream, enabling synchronized analysis of trading activity and order flow dynamics.

4.2 Event constructor




Figure 2: Event construction for 4 order types through 3 questions. Each question is a decision tree for classification. We start with question 1 to question 3. Note, market order is a case of liquidity decreasing.

To apply point process modeling to high-frequency limit order book (LOB) data, we represent market dynamics as a sequence of **event sets**. We consider the top $K = 10$ levels on each side (bid and ask).

Atomic Events. We define a finite set E of *atomic event types*, each denoted $e = (c, s, \delta)$, where:

- $c \in \{\text{LO}, \text{MO}, \text{CO}, \text{IS}\}$ is the *event class*:
 - **Limit Order (LO)**: liquidity provision at one of the top- K levels.
 - **Cancel Order (CO)**: liquidity withdrawal at one of the top- K levels.
 - **Market Order (MO)**: liquidity consumption at the best quote.
 - **Inside-Spread Order (IS)**: a new order placed inside the spread, improving the prevailing best quote.
- $s \in \{b, a\}$ is the *side* (bid or ask).
- δ is the *level index*:
 - For LO/CO, $\delta \in \{1, \dots, K\}$ indicates distance from the best quote ($\delta = 1 = \text{best bid/ask}$, $\delta = 10 = \text{tenth level}$).
 - For MO, $\delta = 1$ always, since market orders execute against the current best quote.
 - For IS, $\delta \in \{1, \dots, K\}$ denotes the number of ticks by which the new quote improves the spread (e.g., $\text{IS}_{b,2}$ = new bid order two ticks inside the spread).

Event Construction from LOB and Trade Data. As illustrated in Figure 2, each atomic order event is inferred by differencing consecutive Level-2 order book snapshots and aligning them with trade records at identical timestamps. Because trades and book updates are perfectly synchronized in both time and volume, every observed change can be classified without ambiguity through a three-step decision tree. First, if the best-quote liquidity decreases and a trade is recorded at the same price, the change corresponds to a **Market Order**—a direct execution against standing quotes. Second, if liquidity decreases at any level without a matching trade, the reduction represents a **Cancel Order**, reflecting liquidity withdrawal. Finally, when liquidity increases, the classification depends on the price position: if the new order appears inside the spread, it forms an **Inside-Spread Order**; otherwise, it is a **Limit Order** placed within the existing book depth. This rule set ensures exhaustive mapping from raw order book and trade data into atomic order

events, enabling a consistent foundation for downstream temporal modeling and anomaly detection.

Event Sets. In crypto limit order books, each event is associated with a set of items. In other words, multiple atomic events may occur simultaneously at the same timestamp. We therefore allow *event sets* at each time t_i :

$$X_{t_i} = \{(e_{i_1}, v(e_{i_1})), \dots, (e_{i_n}, v(e_{i_n}))\},$$

where $e_{i_n} \in \mathcal{E}$ are atomic events and $v(e_{i_n})$ their volume marks.

4.3 State-Dependent Set-MTPP Backbone

Intuition. The event flow of a limit order book is governed not only by past order-flow history, but also by the contemporaneous market condition. In particular, the relative dominance of bid-side and ask-side queues affects the likelihood of future submissions, cancellations, executions, and inside-spread improvements. Following the state-dependent neural Hawkes formulation of Shi and Cartledge, we incorporate a discrete market state into the intensity model so that TFOW captures both temporal dependence and the current order-book regime. Unlike Shi and Cartledge, however, we do *not* adopt the parallel CT-LSTM architecture. We only use their state mechanism and retain a single continuous-time neural backbone, which is more suitable for our set-valued marked event representation and keeps the model compact.

Market state. Let $v_b^{(1)}(t)$ and $v_a^{(1)}(t)$ denote the top-of-book bid and ask queue volumes at time t . We define the queue-imbalance indicator

$$I(t) = \frac{v_b^{(1)}(t) - v_a^{(1)}(t)}{v_b^{(1)}(t) + v_a^{(1)}(t)}.$$

Given a threshold $\theta \in (0, 1)$, we discretize the market state as

$$s_t = \begin{cases} 0, & I(t) \in [-1, -\theta), \\ 1, & I(t) \in [-\theta, \theta], \\ 2, & I(t) \in (\theta, 1]. \end{cases}$$

These three regimes correspond to ask-dominant, balanced, and bid-dominant markets, respectively.

State-dependent hidden dynamics. Let $\mathcal{H}_t$ denote the history of all event sets and volume marks up to time t . As defined in Section 4.2, each timestamp t_i is associated with an event set

$$X_{t_i} = \{(e_{i_1}, v(e_{i_1})), \dots, (e_{i_n}, v(e_{i_n}))\},$$

where $e_{im} \in \mathcal{E}$ is an atomic event and $v(e_{im})$ is its volume mark. TFOW encodes the sequence of past event sets together with the current market state into a continuous-time hidden representation h_t . Between event times, the hidden state decays continuously; at each event time, it is updated using the newly observed event set and the contemporaneous state. This follows the state-dependent neural point process principle that the intensity is decoded from a latent state rather than specified by an additive Hawkes kernel.

We define the baseline intensity by

$$\lambda_\theta^{\text{base}}(t \mid \mathcal{H}_t, s_t) = f_\theta(h_t, s_t),$$

where $f_\theta(\cdot)$ is a positive decoder. In practice, f_θ may be implemented using a softplus output layer to ensure non-negativity of the intensity.

State transition. The market state is allowed to evolve jointly with the event process. Following Shi and Cartledge, after the event set at time t_i is observed, the next state is governed by an event-conditioned transition distribution

$$P(s_{t_i} \mid \mathcal{H}_{t_i}^-, y_i, s_{t_i}^-) \sim \phi_{y_i, s_{t_i}^-},$$

where y_i is the event-type summary associated with timestamp t_i , and $\phi \in \mathbb{R}^{K \times W \times W}$ is a discrete transition tensor. In our implementation, ϕ is estimated offline via empirical transition counting, separately from the neural network parameters.

Compensator. The compensator associated with the baseline state-dependent intensity is

$$\Lambda_\theta^{\text{base}}(t) = \int_0^t \lambda_\theta^{\text{base}}(u \mid \mathcal{H}_u, s_u) du.$$

This baseline Set-MTPP models the structural dynamics of normal order flow. The next subsection introduces a second-order diagnostic derived from compensator increments and uses it to adapt the baseline intensity online.

4.4 Differential 3S as a Second-Order Diagnostic

Intuition. A correctly specified temporal point process becomes a unit-rate Poisson process after the random time change induced by its compensator. Therefore, on the transformed clock, the inter-event gaps should behave like i.i.d. $\text{Exp}(1)$ random variables. This implies not only that the model gets the expected event count right, but also that it gets the geometry of event spacing right. In particular, if $\Delta u_i \sim \text{Exp}(1)$, then

$$\mathbb{E}[\Delta u_i] = 1, \quad \text{Var}(\Delta u_i) = 1, \quad \mathbb{E}[(\Delta u_i)^2] = 2.$$

The original 3S statistic is built from squared transformed gaps and therefore probes this second-moment structure. Intuitively, first-order residuals ask whether the model gets the *number of arrivals* right, while squared compensator increments ask whether it gets the *irregularity of event timing* right. This motivates a sequential second-order diagnostic for TFOW.

Compensator increments. Let

$$u_i = \Lambda_\theta^{\text{base}}(t_i), \quad \Delta u_i = u_i - u_{i-1} = \int_{t_{i-1}}^{t_i} \lambda_\theta^{\text{base}}(u \mid \mathcal{H}_u, s_u) du.$$

If the baseline state-dependent model is correctly specified, then by the random time-change theorem,

$$\Delta u_i \stackrel{i.i.d.}{\sim} \text{Exp}(1).$$

Hence,

$$\mathbb{E}[(\Delta u_i)^2] = 2.$$

This second moment is the key quantity behind 3S. Rather than computing a global goodness-of-fit statistic over an entire transformed sequence, we extract an event-level residual from the same quadratic intuition.

Differential 3S residual. We define the differential 3S residual by

$$r_i^{(2)} = (\Delta u_i)^2 - 2.$$

Under the correctly specified model,

$$\mathbb{E}[r_i^{(2)}] = 0.$$

Thus, $r_i^{(2)}$ measures the local second-order deviation of transformed spacing from the Poisson benchmark. Positive values indicate abnormally irregular, bursty, or dispersed spacing on the compensator clock, whereas negative values indicate overly regular spacing. In this sense, $r_i^{(2)}$ is a quadratic second-order diagnostic on compensator increments.

Smoothed deviation state. To obtain a stable sequential signal, we smooth the event-level residual recursively:

$$d_i^{(2)} = \rho d_{i-1}^{(2)} + (1 - \rho) r_i^{(2)}, \quad \rho \in [0, 1].$$

The process $d_i^{(2)}$ summarizes recent second-order deviations in transformed event spacing. Unlike a rolling-window goodness-of-fit test, this construction is incremental and naturally compatible with sequential modeling.

Deviation-controlled intensity. We incorporate the second-order deviation signal into TFOW through a multiplicative correction of the baseline state-dependent intensity:

$$\lambda_\theta(t \mid \mathcal{H}_t, s_t) = \lambda_\theta^{\text{base}}(t \mid \mathcal{H}_t, s_t) \exp(\beta d_i^{(2)}), \quad t \in (t_i, t_{i+1}],$$

where $\beta \in \mathbb{R}$ is learnable. Equivalently,

$$\log \lambda_\theta(t \mid \mathcal{H}_t, s_t) = \log \lambda_\theta^{\text{base}}(t \mid \mathcal{H}_t, s_t) + \beta d_i^{(2)}.$$

This decomposition separates structural order-flow dynamics from second-order temporal irregularity: the baseline term captures the dynamics explained by event history and imbalance state, while the differential 3S term captures residual spacing abnormalities not explained by the structural model.

Discussion. The proposed diagnostic is complementary to first-order residual monitoring. Standard martingale residuals compare realized and expected counts and are therefore sensitive to count drift. In contrast, differential 3S probes the second moment of transformed waiting times and is sensitive to clustering, dispersion, and abnormal spacing geometry. In crypto limit order books, this allows TFOW to respond not only to unusual event volume, but also to strategically timed bursts, pauses, and irregular cancellation or inside-spread patterns that are characteristic of manipulative flow.

4.5 Differential 3S as a Second-Order Diagnostic on Compensator Increments

4.5.1 State-Dependent Intensity Modeling. In limit order book markets, the arrival of events is influenced not only by past order flow but also by the contemporaneous market condition. Recent work therefore extends neural Hawkes processes by conditioning the event intensity on a discrete market state that summarizes the current order book configuration [12].

Formally, let H_t denote the event history and s_t a state variable derived from the order book. The conditional intensity for event

type k is then modeled as

$$\lambda_k(t \mid H_t, s_t) = f_k(h_t, s_t), \quad (1)$$

where h_t is a latent representation of past events produced by a continuous-time recurrent mechanism and $f_k(\cdot)$ decodes the instantaneous arrival rate. The state variable typically partitions the market into regimes such as bid-dominant, balanced, or ask-dominant based on queue imbalance.

We extend the intensity formulation by introducing a deviation-controlled feedback mechanism that enables explicit separation between structural dynamics and market-driven temporal irregularity.

Baseline TFOW intensity. Model's conditional intensity is

$$\lambda_0(t) = \lambda(h_t, x_t), \quad (2)$$

where h_t encodes historical order flow and x_t denotes the imbalance regime state. The associated compensator is

$$\Lambda_0(t) = \int_0^t \lambda_0(\tau) d\tau. \quad (3)$$

The derivation of $\lambda_0(t)$ and $\Lambda_0(t)$ follows the temporal point process.

$$\Delta u_i = \Lambda(t_i) - \Lambda(t_{i-1}) = \int_{t_{i-1}}^{t_i} \lambda(\tau) d\tau. \quad (4)$$

Under correct model specification, the time-rescaling theorem implies

$$\Delta u_i \sim \text{Exp}(1),$$

i.e., rescaled inter-arrival times are i.i.d. exponential with unit mean.

Intuition. A correctly specified temporal point process should become a unit-rate Poisson process after the random time change induced by its compensator. On this transformed clock, the inter-event gaps should behave like i.i.d. $\text{Exp}(1)$ variables. This implies not only a correct *first-order* behavior, i.e., the expected number of arrivals matches the compensator, but also a correct *second-order* behavior, i.e., the transformed waiting times have the right dispersion and spacing geometry. In particular, if $\Delta u_i \sim \text{Exp}(1)$, then

$$\mathbb{E}[\Delta u_i] = 1, \quad \text{Var}(\Delta u_i) = 1, \quad \mathbb{E}[(\Delta u_i)^2] = 2.$$

The classical 3S statistic is built from squared transformed gaps and therefore probes this second-moment structure. Intuitively, while first-order residuals measure whether the model gets the *event count* right, squared compensator increments measure whether the model gets the *timing irregularity* right. Large values indicate abnormally bursty or dispersed spacing on the transformed clock, while small values indicate overly regular spacing. This makes the second moment of compensator increments a natural second-order diagnostic for local model misspecification.

Formally, let $\lambda_\theta(t \mid \mathcal{H}_t)$ denote the conditional intensity produced by TFOW, and define its compensator

$$\Lambda_\theta(t) = \int_0^t \lambda_\theta(u \mid \mathcal{H}_u) du.$$

For the observed event times $\{t_i\}_{i=1}^N$, define the transformed times

$$u_i = \Lambda_\theta(t_i), \quad \Delta u_i = u_i - u_{i-1}.$$

By the random time-change theorem, if the model is correctly specified, then the transformed process is a standard Poisson process and the increments satisfy

$$\Delta u_i \stackrel{i.i.d.}{\sim} \text{Exp}(1).$$

Motivated by the quadratic structure of 3S, we define the event-level differential 3S residual

$$r_i^{(2)} = (\Delta u_i)^2 - 2.$$

Since $\mathbb{E}[(\Delta u_i)^2] = 2$ under the Poisson null, $r_i^{(2)}$ is centered:

$$\mathbb{E}[r_i^{(2)}] = 0.$$

Thus, $r_i^{(2)}$ quantifies the local second-order deviation of the transformed spacing from the model-implied $\text{Exp}(1)$ benchmark. Specifically, $r_i^{(2)} > 0$ indicates abnormally irregular transformed spacing, whereas $r_i^{(2)} < 0$ indicates overly regular spacing.

To obtain a stable sequential signal, we smooth the event-level residual recursively:

$$d_i^{(2)} = \rho d_{i-1}^{(2)} + (1 - \rho) r_i^{(2)}, \quad \rho \in [0, 1].$$

The state $d_i^{(2)}$ summarizes recent second-order deviations on the compensator clock. We then incorporate it into the conditional intensity as a multiplicative correction:

$$\lambda_\theta(t \mid \mathcal{H}_t) = \lambda_\theta^{\text{base}}(t \mid \mathcal{H}_t) \exp(\beta d_i^{(2)}), \quad t \in (t_i, t_{i+1}],$$

where $\lambda_\theta^{\text{base}}$ is the base TFOW intensity and β controls the sensitivity to second-order deviation.

This construction is closely related in spirit to 3S, since both are driven by squared transformed inter-event gaps. The distinction is that the original 3S statistic is a *global* goodness-of-fit functional, whereas our formulation extracts an *incremental* second-order residual process suitable for sequential modeling. Consequently, TFOW is able to respond online not only to the structural order-flow state, but also to recent deviations in the timing geometry of events relative to the compensator-implied Poisson benchmark.

Discussion. The proposed diagnostic is complementary to first-order residual monitoring. Standard martingale residuals compare observed and expected counts through $N(t) - \Lambda_\theta(t)$, and are therefore sensitive to count drift. In contrast, $r_i^{(2)}$ probes the second moment of transformed waiting times and is sensitive to local irregularity in event spacing. In microstructure settings, this allows the model to detect regime shifts that manifest not only through excess event volume, but also through abnormal clustering, dispersion, or regularization of order-flow timing.

4.6 State-Dependent Intensity with Deviation Feedback

Deviation-controlled intensity. We define the modified intensity:

$$\lambda(t) = \lambda_0(t) \cdot \exp(\beta s_t), \quad (5)$$

with $\beta \in \mathbb{R}$ learnable.

Taking logarithms,

$$\log \lambda(t) = \log \lambda_0(t) + \beta s_t. \quad (6)$$

The baseline term $\lambda_0(t)$ captures structural dynamics encoded in (h_t, x_t) , which summarize event history and regime information. Any deviation in arrival behavior not explained by these structural components must manifest through systematic distortions of rescaled-time spacings.

Since s_t is computed solely from compensator increments $\{\Delta u_i\}$, it measures inconsistency between predicted and realized temporal spacing. Importantly, s_t depends only on the temporal arrangement of events relative to the model's own predicted intensity and under correct structural modeling, fluctuations in s_t arise from genuine clustering or dispersion in the market. Persistent elevation of s_t corresponds to bursty arrivals (volatility shocks, liquidation cascades), and suppressed s_t corresponds to under-dispersed flow (liquidity drought, inactivity).

Therefore, βs_t quantifies multiplicative amplification (or attenuation) of structural intensity required to reconcile predicted flow with observed temporal irregularity.

In other words,

$$\beta s_t = \log \frac{\lambda(t)}{\lambda_0(t)}, \quad (7)$$

which measures the deviation of realized market activity from structurally expected activity.

Let $\hat{\lambda}_{\text{obs}}(t)$ denote empirical intensity estimated over a short window. We define

$$D_{\text{market}}(t) = \beta s_t, \quad (8)$$

$$D_{\text{model}}(t) = \left| \log \hat{\lambda}_{\text{obs}}(t) - \log \lambda(t) \right|. \quad (9)$$

Then

$$\log \hat{\lambda}_{\text{obs}}(t) = \log \lambda_0(t) + D_{\text{market}}(t) + \epsilon_t, \quad (10)$$

where ϵ_t corresponds to residual model misspecification.

Interpretational summary.

- $\lambda_0(t)$: structural prediction from TFOW history and regime.
- βs_t : systematic temporal irregularity induced by market dynamics.
- ϵ_t : unexplained residual error.

Thus, βs_t represents market-driven deviation because it quantifies the multiplicative adjustment necessary to align structural intensity with empirically observed clustering, independently of latent state encoding.

This decomposition renders temporal abnormality identifiable and enables real-time attribution of intensity variation to market shocks rather than model error.

4.6.1 Model Training. !!!* train baseline intensity first * compute differential 3S from detached compensator increments * feed it back as a correction feature

Modeling Event Type and Volume in Order book Data. The event constructor outputs a sequence of set of events. Formally, let each event at time t_i be represented as a tuple

$$x_i = (e_i, v_i),$$

where e is the event type and $v_i \in \mathbb{R}^+$ is the corresponding volume. We factorize the conditional mark distribution as

$$p(x_i \mid t_i, H(t_i)) = p(e_i \mid t_i, H(t_i)) \cdot p(v_i \mid e_i, t_i, H(t_i)), \quad (11)$$

where $H(t_i)$ denotes the event history up to time t_i .

Set-Valued Marked Temporal Point Processes (Set-MTPP). Classical marked temporal point processes (MTPPs) assume that each event is associated with a single categorical mark. This simplification breaks down in crypto limit order book level-2 data nature, where multiple atomic events—such as simultaneous limit placements, cancellations, and executions—often occur at exactly the same timestamp. To properly capture these co-occurrences, we adopt the *set-valued MTPP* framework [1], which generalizes neural MTPPs from singleton events to event sets.

Formally, let the event history be

$$H(T) = \{(\tau_i, X_i)\}_{i=1}^N,$$

where $\tau_i \in \mathbb{R}^+$ is the i -th event time and $X_i \subseteq \{1, \dots, M\}$ is a finite subset of atomic event type e . The history up to time t is $H(t) = \{(\tau_j, X_j) : \tau_j < t\}$.

The process is parameterized by a ground intensity $\lambda^*(t)$ and a conditional distribution over sets $p^*(X | t, H(t))$. The marked intensity factorizes as

$$\lambda_X^*(t) = \lambda^*(t) p^*(X | t, H(t)).$$

Accordingly, the joint log-likelihood decomposes into temporal and set contributions:

$$\mathcal{L}(H(T)) = \underbrace{\sum_{i=1}^N \log \lambda^*(\tau_i) - \int_0^T \lambda^*(s) ds}_{\mathcal{L}_{\text{time}}} + \underbrace{\sum_{i=1}^N \log p^*(X_i | \tau_i, H(\tau_i))}_{\mathcal{L}_{\text{set}}}.$$

Set-MTPPs therefore provide a principled generative framework for level-2 LOB modeling, separating the dynamics of *when* events occur from *which sets* of events occur. This enables likelihood-based inference over simultaneous market microstructural patterns.

Volume Set Multivariate Temporal Point Processes. While set-MTPPs capture simultaneous event types, each LOB event also carries a continuous *volume mark* that reflects liquidity provision or consumption. Volumes are heavy-tailed and constrained by available depth—features that categorical marks alone cannot capture. We therefore extend set-MTPPs to a *volume-set-MTPP*, which models both the composition of event sets and their associated traded volumes under realistic market constraints.

At each event time t_i , we observe:

$$X_i \subseteq \{1, \dots, M\}, \quad \mathbf{v}_{t_i} = \{v(e) \in \mathbb{R}_{>0} : e \in X_{t_i}\},$$

where X_{t_i} is the set of event types and $\mathbf{v}_{t_i}$ their associated volumes. The modeling task is to learn

$$p(X_t, \mathbf{v}_t | H_t),$$

where H_t is the market history up to t .

Set-Volume Factorization. We decouple set composition and volume magnitudes:

$$p(X_t, \mathbf{v}_t | H_t) = p(X_t | H_t) \prod_{e \in X_t} p(v(e) | e, H_t, \text{cap}_t(e)).$$

This allows the model to (i) capture co-occurrence structure in $p(X_t | H_t)$, (ii) model event-specific volume distributions, and (iii) respect market microstructure constraints via $\text{cap}_t(e)$.

Dynamic Set Model. Event sets are modeled via conditionally independent Bernoulli variables:

$$p(X_t = x | H_t) = \prod_{k=1}^K \rho_k(h_t)^{x_k} (1 - \rho_k(h_t))^{1-x_k},$$

with $\rho_k(h_t) = \sigma(\mathbf{w}_k^\top \mathbf{n}(h_t) + b_k)$ parameterized by the hidden state h_t .

Neural Architecture and Training. We employ a Neural Hawkes [9] to produce hidden states $h_{d,t}$, with three prediction heads: (i) an intensity head for timing $\lambda(t)$, (ii) a set head for $\{\rho_k(h_{d,t})\}$, and (iii) a volume head for $\{(\mu_k(h_{d,t}), \sigma_k(h_{d,t}))\}$. The loss decomposes as

$$\mathcal{L} = \mathcal{L}_{\text{timing}} + \mathcal{L}_{\text{sets}} + \mathcal{L}_{\text{volumes}},$$

balancing temporal, set, and volume likelihood terms.

This formulation yields a principled and high-fidelity generative model of LOB microstructure—enabling both simulation and anomaly detection in high-frequency trading.

4.7 Anomaly Detection

Market manipulators deliberately inject crafted orders into otherwise normal order flows, perturbing the natural distribution of event timings, traded volumes, and directional dependencies between orders. Such distortions are inherently statistical and can therefore be detected through distributional tests on the reconstructed event history.

Formally, let the full observed history up to horizon T be

$$\mathcal{H}(T) = \{(\tau_i, X_{t_i})\}_{i=1}^N,$$

where $t_1 < t_2 < \dots < t_N \leq T$ are event times and X_{t_i} denotes the event set observed at time t_i .

Stage 1: Sequence-Level Detection. We first segment the full history into hourly subsequences and test each for deviations from the learned market baseline. Specifically, we define

$$\mathcal{H}_i = \mathcal{H}[t_a, t_b] = \{(\tau_j, X_{t_j}) : t_a < \tau_j \leq t_b\}, \quad D_{\text{train}} = \{\mathcal{H}_1, \dots, \mathcal{H}_M\},$$

where the subsequences in D_{train} are assumed to follow the empirical distribution P_{data} . For every $\mathcal{H}_i$, we conduct the hypothesis test

$$H_0 : \mathcal{H}_i \sim P_{\text{data}} \quad \text{vs.} \quad H_1 : \mathcal{H}_i \sim Q, \quad Q \neq P_{\text{data}}.$$

To quantify deviations, we apply the *sum-of-squared-spacings* (3S) statistic:

$$\psi(Z) = \frac{1}{V} \sum_{i=1}^{N+1} (v_i - v_{i-1})^2, \quad Z = (\Lambda^*(t_1), \dots, \Lambda^*(t_N)),$$

where $\Lambda^*(t)$ is the fitted compensator of the model intensity and $Z = \Lambda^*(T)$ the total rescaled horizon. The 3S statistic simultaneously captures irregularities in event spacing and cumulative mass flow. Subsequences with low p -values under P_{data} are flagged as anomalous trading hours.

Stage 2: Event-Level Localization. For each detected anomalous window $\mathcal{H}[a_t, b_t]$, TFOW proceeds to *localize* the responsible atomic events through probabilistic queries derived from the fitted process.

(a) Timestamp-level test. We first evaluate *hitting-time anomalies*—rare early arrivals of specific event types—by computing

$$\mathbb{P}(\text{hit}(A) < t \mid \mathcal{F}_{t_a}), \quad \mathcal{F}_{t_a} = \mathcal{H}[0, t_a] = \{(t_j, X_{t_j}) : t_j < t_a\}.$$

Significantly smaller values than expected under P_{data} indicate abnormally early occurrences of events in subset A , pointing to potential manipulative bursts.

(b) Motif-level test. When such deviations are present, we escalate to *motif-level localization*, focusing on the minimal temporal motif between two consecutive timestamps $(t_{i-1}, t_i]$. For disjoint event subsets A, B , we compute the directional probability

$$q(A < B) = \Pr(t_A < t_B \mid P_{\text{data}}),$$

which measures how often A precedes B under the learned dynamics. Anomalous small $q(A < B)$ values expose rare directional orderings—interpretable microstructural signatures of manipulative flow.

This two-stage design ensures a consistent “small-is-anomalous” interpretation: the 3S test surfaces statistically irregular trading intervals; hitting-time queries highlight unexpectedly early local arrivals; and motif queries explain directional distortions between consecutive event sets. Together they yield both *statistical validity* and *microstructural interpretability*, enabling TFOW to pinpoint manipulative behaviors at order-level granularity.

4.7.1 Anomaly Sequence Detection. We extend the probabilistic testing paradigm of Shchur et al. [10] to the limit order book setting with set-valued events and continuous volumes. The central idea is to map observed event sequences into the standard Poisson process (SPP) via appropriate compensators, and then apply *sum-of-squared-spacings* (3S) tests. This unifies the detection of anomalies in timing, spacing, and traded volumes.

Sum-of-Squared-Spacings. We adopt the sum-of-squared-spacings (3S) statistic for order book anomaly detection. Given fitted total intensity

$$\lambda^*(t) = \sum_{x \in \mathcal{X}} \lambda_x^*(t)$$

, the compensator

$$\Lambda^*(t) = \sum_{x \in \mathcal{X}} \Lambda_x^*(t) = \sum_{x \in \mathcal{X}} \int_0^t \lambda_x^*(u) du$$

maps arrivals t_i to rescaled times $u_i = \Lambda^*(t_i)$. The statistic

$$\psi = \frac{1}{\Lambda^*(T)} \sum_{i=1}^{N+1} (u_i - u_{i-1})^2$$

quantifies spacing irregularities and count deviations simultaneously. Unlike KS-based tests, 3S captures bursts, lulls, and atypical event counts; when extended with volume compensators, it further detects abnormal liquidity flows, making it well suited for identifying manipulative order flow.

Mass-Flow Deviations. Manipulation often manifests not in counts but in *abnormal volume flow*. We define the *mass intensity*

$$\rho^*(t) = \sum_{k=1}^K \lambda_k^*(t) \mathbb{E}[v \mid k, t, H_t], \quad \Xi^*(t) = \int_0^t \rho^*(u) du,$$

where $\mathbb{E}[v \mid k, t, H_t]$ is the truncated log-normal expectation from our volume head. Mapping arrivals to $m_i = \Xi^*(t_i)$ yields the mass-based statistic

$$\psi_{\text{mass}} = \frac{1}{\Xi^*(T)} \sum_{i=1}^{N+1} (m_i - m_{i-1})^2,$$

which flags anomalies where mass accumulates atypically.

4.7.2 Anomaly Localization. To achieve both statistical rigor and interpretability, we employ a two-stage anomaly localization. First, the joint 3S detector is applied to *hourly subsequences*, producing anomaly signals at the scale of trading hours. Second, within each anomalous hourly subsequence, we conduct *timestamp-level query*: for every timestamp t_i , we compute the conditional probability of its set given the immediately preceding timestamp under the fitted volume-set-MTPP. If this probability falls below a threshold, α , we escalate to the motif query over $(t_{i-1}, t_i]$. By ranking anomalous motif $A < B$, we retain the top- N most anomalous consecutive sets as the *anomaly motifs*. These motifs reveal interpretable precursors of anomalies.

Hitting Time Queries. Let $A \subseteq \{1, \dots, M\}$ be a non-empty subset of atomic events. The *hitting time* $\text{hit}(A)$ is defined as the first occurrence time τ such that $X_\tau \cap A \neq \emptyset$. The probability that any event in A occurs before a horizon t can be expressed as

$$\mathbb{P}(\text{hit}(A) \leq t) = 1 - \mathbb{E}_{H[0,t] \sim q} \left[\exp \left(- \int_0^t \lambda^*(s) p^*(A \mid s, H(s)) ds \right) \right],$$

where $\lambda^*(t)$ is the ground intensity and $p^*(A \mid t, H(t))$ is the conditional probability that the sampled set at time t contains at least one item from A . For the Dynamic Bernoulli factorization, this reduces to

$$p^*(A \mid t, H(t)) = 1 - \prod_{k \in A} (1 - p^*(k \mid t, H(t))).$$

This formulation allows efficient Monte Carlo estimation under importance sampling and provides a natural way to quantify the likelihood of early occurrences, which we later use for identifying anomalously early event clusters.

A-before-B Queries. For two disjoint subsets $A, B \subseteq \{1, \dots, M\}$, we are interested in the probability that some event in A occurs before any event in B within horizon t , i.e.,

$$\mathbb{P}(\text{hit}(A) < \text{hit}(B), \text{hit}(A) \leq t).$$

Using the factorization $\lambda_X^*(t) = \lambda^*(t) p^*(X \mid t, H(t))$ and importance sampling with proposal intensity restricted to $(A \cup B)$, the unbiased estimator is

$$\begin{aligned} & p(\text{hit}(A) < \text{hit}(B), \text{hit}(A) \leq t) \\ &= \mathbb{E}_{H[0,t]}^q \left[\int_0^t \exp \left(- \int_0^s \lambda_{A \cup B}^*(s') ds' \right) \lambda_{A \cap \neg B}^*(s) ds \right], \quad (12) \end{aligned}$$

where $\lambda_S^*(t) = \lambda^*(t) p^*(S \mid t, H(t))$ for any subset S . The term $p^*(A \cap \neg B \mid t, H(t))$ represents the probability that the sampled set at t contains an item in A but none in B . Symmetric forms yield estimators for the other three cases— B before A , simultaneous occurrence, and neither before t . These directional queries form

the basis for detecting causal or lead-lag relations between distinct event types in continuous time.

Localization Procedure. Formally speaking, given a detected subsequence $[t_0, t_n]$ and history filtration $\{\mathcal{F}_t\}$, we test anomalies in two escalating stages. First, for $i = 1, \dots, n$, conditioning on $\mathcal{F}_{t_{i-1}}$, we compute the local hazard of hitting $A \subseteq \{1, \dots, K\}$ on $(t_{i-1}, t_i]$, with compensator $\Lambda_A^*(t_{i-1}, t_i) = \int_{t_{i-1}}^{t_i} \lambda^*(s) p^*(A | s, H(s)) ds$. The one-sided p -value $p_i^{\text{hit}}(A) = \exp(-\Lambda_A^*(t_{i-1}, t_i))$ is small when A appears earlier than expected, flagging rare local hits. Second, if $p_i^{\text{hit}}(A) < \alpha_1$, we estimate the conditional probability that A precedes disjoint B in $(t_{i-1}, t_i]$:

$$q_i(A \prec B) = \int_{t_{i-1}}^{t_i} \exp\left(-\int_{t_{i-1}}^s \lambda_{A \cup B}^*(u) du\right) \lambda_{A \cap B}^*(s) ds.$$

Here, small $q_i(A \prec B)$ values indicate that the motif $A \prec B$ is rare under the model, and thus anomalous. We therefore label $(t_{i-1}, t_i]$ as anomalous and record motif $(A \prec B)$ iff $p_i^{\text{hit}}(A) < \alpha_1$ and $q_i(A \prec B) < \alpha_2$. This construction ensures a consistent “small-is-anomalous” interpretation across both stages, yielding motifs that are simultaneously surprising in timing and rare in directional ordering.

This framework provides a statistically principled path from *detection* to *localization* to *explanation*. The 3S statistic unifies timing, spacing, and mass anomalies, but it cannot be reliably applied to sequences with very few timestamps, since the variance of the spacing distribution is ill-defined in such regimes. Our query-based refinement overcomes this limitation by leveraging probabilistic set-MTPP queries, ensuring robustness in sparse cases while extracting interpretable temporal motifs. Together, the framework delivers a high-precision, label-free anomaly detector tailored to the structure of high-frequency financial markets with real-world interpretability.

5 EVALUATION

5.1 Evaluation Setup

5.2 Evaluation Results

1. tsne 62 types but with 4 types colored
2. table comparison of log likelihood, rsme, error rate, —Transformer Hawkes Process
- a. hawkes process b.LSTM c. neural hawkes d. SAHP (self-attentive neural Hawkes process) e. state-dependent neural Hawkes
3. kstylist fashion of sim vs real
4. set vs no set (accuracy)

6 RELATED

temporal point process anomaly detection

Goodness of fit in neural network Residual TPP: A Unified Lightweight Approach for Event Stream Modeling (ICLR / OpenReview 2025)

Neural Networks with Recurrent Generative Feedback self-referential structure adaptive filter predictive coding boosting

7 DISCUSSION AND CONCLUSION

8 ACKNOWLEDGMENT

Sincere thanks to Konark Jain and Antonio Briola’s guidance on order book microstructures.

REFERENCES

- [1] Yuxin Chang, Alex Boyd, and Padhraic Smyth. 2024. Probabilistic Modeling for Sequences of Sets in Continuous-Time. In *International Conference on Artificial Intelligence and Statistics*, Vol. 238. 4357–4365.
- [2] Cloudburst. 2025. Cloudburst Technologies. <https://burst.cloud/>
- [3] Lin William Cong, Xi Li, Ke Tang, and Yang Yang. 2023. Crypto Wash Trading. *Management Science* 69, 11 (2023), 6427–6454. <https://doi.org/10.1287/MNSC.2021.02709>
- [4] Sumitra Ganesh, Nelson Vadori, Jpmorgan Ai, Research Mengda, Xu Jpmorgan, A I Research, Hua Zheng, Prashant Reddy, and Manuela Veloso. [n. d.]. Reinforcement Learning for Market Making in a Multi-agent Dealer Market. ([n. d.]).
- [5] Sascha Hägele. 2024. Centralized exchanges vs. decentralized exchanges in cryptocurrency markets: A systematic literature review. *Electronic Markets* 34, 1 (12 2024), 33. <https://doi.org/10.1007/s12525-024-00714-2>
- [6] Konark Jain, Nick Firoozee, Jonathan Kochems, and Philip Treleaven. 2024. Limit Order Book dynamics and order size modelling using Compound Hawkes Process. *Finance Research Letters* 69 (11 2024), 106157. <https://doi.org/10.1016/J.FRL.2024.106157>
- [7] Konark Jain, Jean-François Muzy, Jonathan Kochems, and Emmanuel Bacry. 2025. No Tick-Size Too Small: A General Method for Modelling Small Tick Limit Order Books. (8 2025). <http://arxiv.org/abs/2410.08744>
- [8] Kaiko. 2025. The crypto industry’s leading data, analytics, and indices provider. - Kaiko. <https://www.kaiko.com/>
- [9] Hongyuan Mei and Jason M Eisner. 2017. The Neural Hawkes Process: A Neurally Self-Modulating Multivariate Point Process. In *Advances in Neural Information Processing Systems*, I Guyon, U Von Luxburg, S Bengio, H Wallach, R Fergus, S Vishwanathan, and R Garnett (Eds.), Vol. 30. Curran Associates, Inc. https://proceedings.neurips.cc/paper_files/paper/2017/file/6463c88460bd63bbe256e495c63aa40b-Paper.pdf
- [10] Oleksandr Shchur, Ali Caner Turkmen, Tim Januschowski, Jan Gasthaus, and Stephan Günnemann. 2021. Detecting Anomalous Event Sequences with Temporal Point Processes. In *Advances in Neural Information Processing Systems*, M Ranzato, A Beygelzimer, Y Dauphin, P S Liang, and J Wortman Vaughan (Eds.), Vol. 34. Curran Associates, Inc., 13419–13431. https://proceedings.neurips.cc/paper_files/paper/2021/file/6faa8040da20ef399b63a72d0e4ab575-Paper.pdf
- [11] Oleksandr Shchur, Ali Caner Turkmen, Tim Januschowski, and Stephan Günnemann. 2021. Neural Temporal Point Processes: A Review. *IJCAI International Joint Conference on Artificial Intelligence* (4 2021), 4585–4593. <https://doi.org/10.24963/ijcai.2021/623>
- [12] Zijian Shi and John Cartlidge. 2022. State Dependent Parallel Neural Hawkes Process for Limit Order Book Event Stream Prediction and Simulation. In *Proceedings of the 28th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*. ACM, New York, NY, USA, 1607–1615. <https://doi.org/10.1145/3534678.3539462>
- [13] Utkarsh Upadhyay, Abir De, and Manuel Gomez Rodriguez. 2018. Deep Reinforcement Learning of Marked Temporal Point Processes. *Advances in Neural Information Processing Systems* 31 (2018). <https://github.com/Networks-Learning/tpprl>
- [14] Jiahua Xu and Benjamin Livshits. 2019. *The Anatomy of a Cryptocurrency [Pump-and-Dump] Scheme*. 1609–1625 pages. www.usenix.org/conference/usenixsecurity19/presentation/xu-jiahua